



Ashwanth Kumar

PRINCIPAL ENGINEER

☎ (+91) 96000 87935 | ✉ ashwanthkumar@googlemail.com | 🌐 ashwanthkumar.in | 📱 [ashwanthkumar](#)

tl;dr

7+ years of experience · Loves to understand, build and operate distributed systems · System Operations · Organizes internal / external hackathons · [Talks](#) at local meetups and conferences · OSS Supporter ·

Experience

Avalara

Chennai, India

SENIOR TECH LEAD

Feb. 2019 - Present

- Working on automating detection of Sales and Use Tax changes across various states within US. Reducing time from 3+ weeks to few minutes across different states of US.

Indix

Chennai, India

PRINCIPAL ENGINEER

Oct. 2015 - Jan. 2019

- Lead Engineer on our internal Data Services Platform to run our Big Data Workloads.
- Tech Lead on Product Matching - Enables Price Comparison use-cases and detecting duplicates with our 2B+ Product Catalog.
- Tech Lead in Data Ingestion Team that crawls, parses and stores >20 million webpages a day.
- Finder - HTML Archive Storage system for storing over 3 billion HTML documents.
- [Suuchi](#) - OpenSource toolkit to build distributed data systems. Built for Finder initially but later adopted to other projects internally as well.
- Saved million dollars a year on AWS by intelligently using Spot instances and Auto Scaling for 70% of our workloads.

Indix

Chennai, India

SOFTWARE ENGINEER

Jun. 2012 - Sept. 2015

- First generation of Indix Data Pipeline which is completely based out of λ architecture.
- Two generations of our web crawler - scaling to millions of pages per day across 1000s of websites.
- Two generations of Scheduler which feeds our crawler with the right urls every day.
- Three generations of Parser. Built a DSL to express transformations in a language agnostic way that scaled the parsers to 1000s of websites.
- From Java based MR jobs to Scalding to building a query tool with a DSL (inspired from SQL and on top of scalding) which enabled engineers to access the data much faster for analysis / debugging.

Mu Sigma

Bangalore, India

BIG DATA INTERN

Jan. 2012 - May. 2012

- Evaluating AWS for migrating to cloud from custom data centers
- Worked on running R code on Storm.

Skills

DevOps Tools AWS, Mesos, Marathon, Docker, Chef, Ansible, Basic Linux Administration

Programming Scala, Java, Go, Node.JS

Big Data Hadoop, HBase, Spark, Zookeeper, Kafka, Samza, RocksDB

Talks

CONFERENCES

April 2018 **Why we built a distributed system**, DSConf 2018 Pune, India [Slides](#) | [Video](#)

Nov. 2017 **Lessons scaling operations to everyone @indix**, Mini Cloud Conf Chennai, India [Slides](#) | [Video](#)

Nov. 2017 **Using Monoids for Large Scale Aggregates**, Scala.io Lyon, France [Slides](#) | [Video](#)

Nov. 2016 **Lessons from managing Hadoop clusters on AWS @indix**, DevopsDays India [Slides](#) | [Video](#)